

Introduction to concurrency

Caveat programmer!

- The final two sections are an introduction to the concurrency facilities offered by C++11
- This is not a course on how to:
 - design, write, debug or test concurrent applications
 - get maximum performance from the machine
- Things that aren't covered include:
 - decomposing the problem domain
 - typical concurrency designs and patterns
 - context switching and grain size
 - the influence of caches
 - mixing I/O and CPU-bound tasks

Concurrency spectrum

- Message passing
 - multi-box: Hadoop, Google MapReduce
 - multi-process: MPI, Erlang
- Shared memory
 - parallel libraries: OpenMP, Threading Building Blocks
 - futures and thread pools: `std::future`, `std::async`, etc
 - atomics: `std::atomic`, lock-free, compare-and-swap
 - raw threads: `std::thread`, `std::mutex`, etc
 - memory fences: `std::memory_order_acquire`, etc

program at the
highest level
you can on this
spectrum!

direct support in C++
+11 libraries for
these approaches

increasing order
of complexity
and effort

Futures and `std::async`

any “callable
object”: function,
lambda or functor

```
// C++11
std::future<int> f = std::async([]{ return long_calc(); });
// do other stuff
int i = f.get();           // return result or block until it's ready
                           // trying to read f again is an error
                           // use std::shared_future<int> instead
```

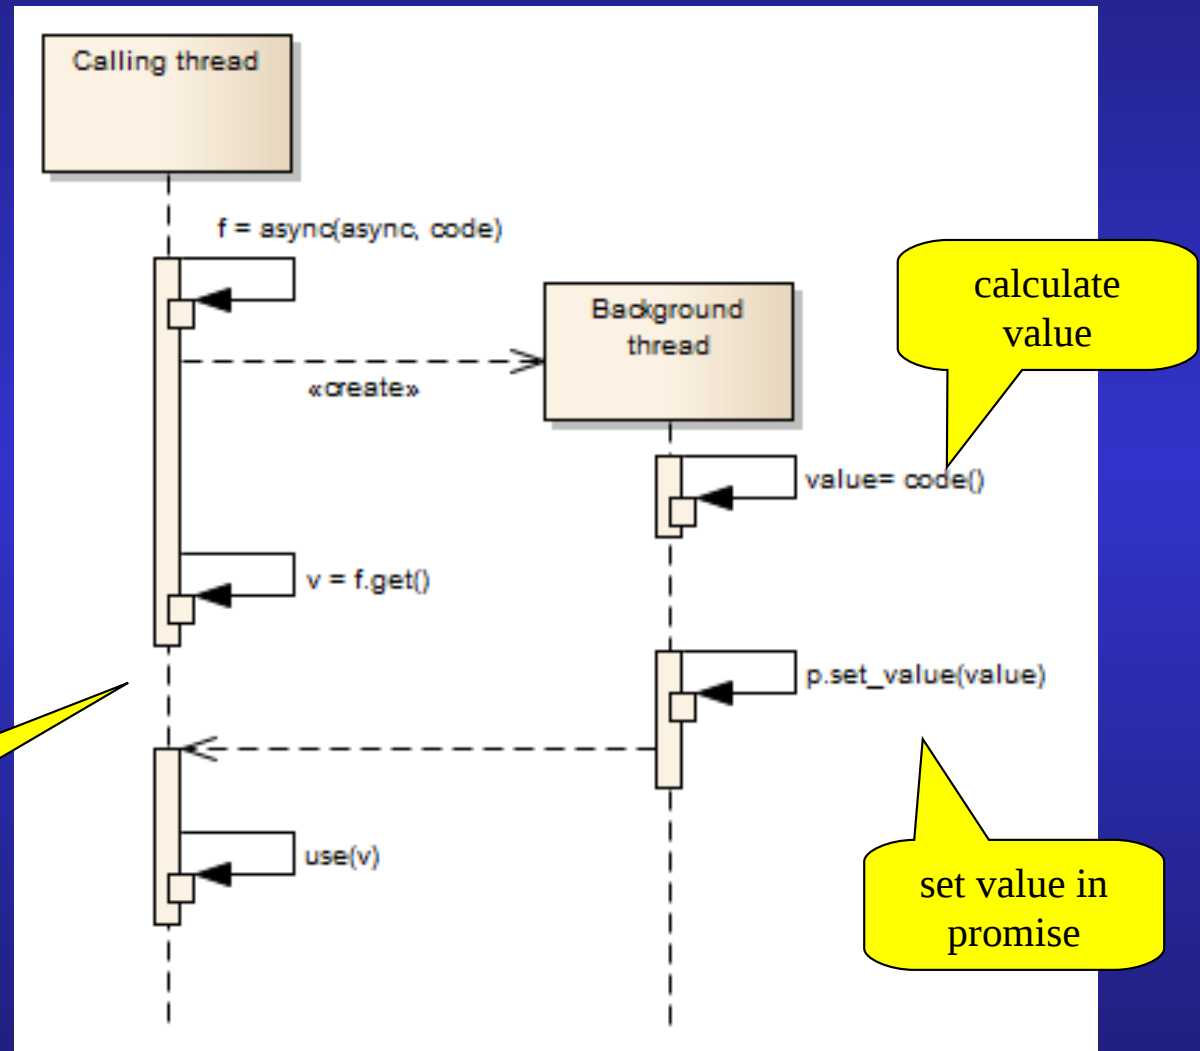
- Futures and promises are two ends of a one-shot message queue
 - promise is the sending end of the queue
 - future returns the value or rethrows exception when `get()` is called
- Destructor of future waits until task is finished
- Async is not a thread pool (likely to appear in TS)

std::async and background thread

```
void thread1()  
{  
    auto f  
        = std::async(  
            std::launch::async,  
            code);  
  
    auto v = f.get();  
    use(v);  
}
```

request
background
thread

block until
promise is set

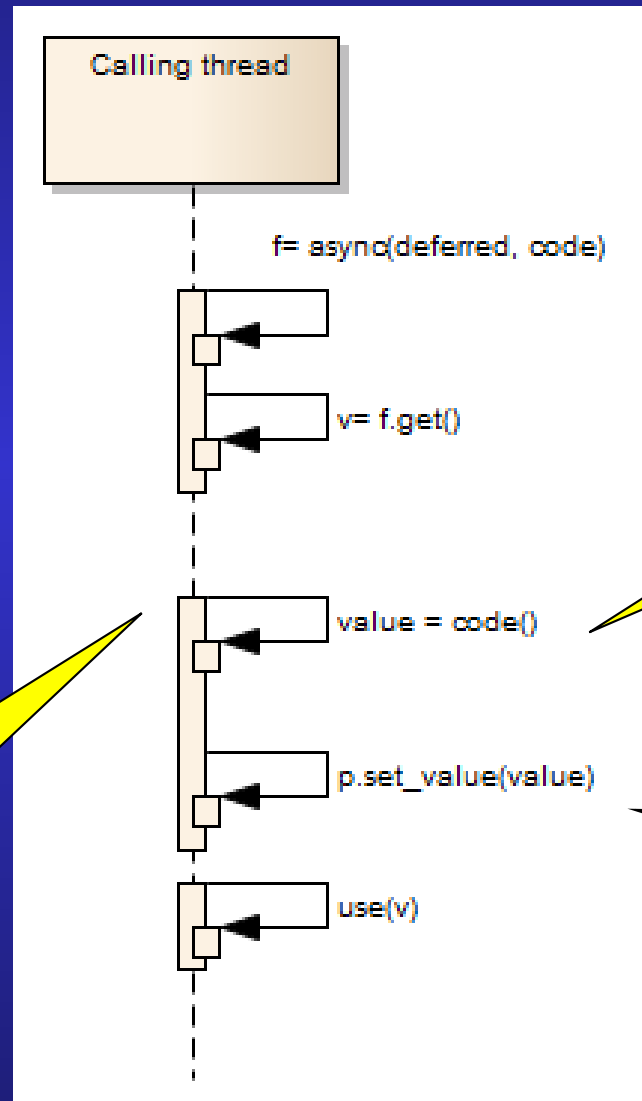


std::async and deferred execution

```
void thread1()  
{  
    auto f  
        = std::async(  
            std::launch::deferred,  
            code);  
  
    auto v = f.get();  
    use(v);  
}
```

request same
thread

calculation
happens on
demand in
same thread



calculate
value

set value in
promise

Four parts to using futures

1. Some code that calculates a value (can be void)
2. A promise that is set by the code
3. A future that waits for the promise to be set
4. A thread that executes the code

- `std::packaged_task` wraps up code with a promise/future pair and is a “callable object”
 - it still needs a thread to run it
- `std::async` with `launch::async` creates a new thread and wraps up the code as a `packaged_task`
- `std::async` with `launch::deferred` wraps up the code as a `packaged_task` and runs it in the current thread

Futures and exceptions

- Normally the promise of a future/promise channel calls `promise.set_value(x)`
- It can also call `promise.set_exception(copy_exception(obj))` and then `future.get()` will rethrow `obj`
 - `make_exception_ptr` instead of `copy_exception`
- Deleting a promise without setting it first will store a `future_error` exception (thrown on `future.get()`)
- `std::exception_ptr` is a reference-counted smart pointer for exceptions

Shared futures

- An ordinary future can be read only once by one thread
 - reading it a second time will throw a `future_error` with a condition of `future_errc::no_state`
- A `std::shared_future` can be shared across threads and can be read multiple times
 - a `shared_future` can be created from a future using a `move` c/tr (requires `std::move(future)`)
- It is **not** safe to access one `shared_future` from multiple threads
- It is safe to access multiple `shared_futures` that access the same `async` result

Concurrency – atomics

- `std::atomic<>` can be applied to built-in types with the usual arithmetic operations (but not `*`, `%` and `/`)
- Also provides load, store, exchange, `compare_exchange_strong/weak` (CAS – lock-free)
- Must be initialised by `c/tr`
- Works with user-defined types with trivial copy
`op= (memcpy) + bitwise == (memcmp)`
- `std::atomic<T>::is_lock_free()`
- Free-function i/f for C compatibility
- Can make use of memory ordering (next section)
 - default is `std::memory_order_seq_cst`

Concurrency – atomics examples

```
std::atomic<int> ai(2);
ai++;                      // usual ops, now ai == 3
std::cout << ai << std::endl;    // implicit load
std::cout << ai.load() << std::endl; // explicit load
auto i = ai.load();

int expect = 4;            // CAS loop - used in lock-free algorithms
while (! ai.compare_exchange_weak(expect, 5))
    std::cout << "i = " << i << ", expect = " << expect << std::endl;

std::cout << "ai = " << ai << std::endl;
// prints: i = 3, expect = 3 // expect gets existing value
//           ai = 5           // update succeeds on second loop
```

- Comp/exchg updates the atomic if the current value is the “expected” value (which is pass-by-ref)
- Otherwise the “expected” value is updated instead
- Strong and weak versions to allow for non-atomic hardware implementations

Further concurrency

Threads

```
// C++11
std::thread t1;                // thread with no associated code
std::thread t2(func);          // run func() with no params
std::thread t3(f, 3, "foo");    // run f(3, "foo")
std::thread t4(&X::m, obj, 5);  // run obj.m(5) (decltype(obj) == X)
std::thread t4a(&X::m, &obj, 5); // run obj->m(5) (obj ptr, not copy)
std::thread t5(g, std::ref(x)); // run g with x passed by reference
std::thread t6(std::move(t4));  // can't copy a thread, only move
                                // t4 is now "empty", like t1

if (t2.joinable()) t2.join();   // threads must be joined or detached
t3.detach();                    // create background thread
```

- Threads must be joined or detached, or an error will occur in the destructor
- `std::thread::hardware_concurrency()` returns the number of available cores

The current thread

```
// C++11
// delays execution of this thread for a given time (or more)
std::chrono::milliseconds fiveMillis(5);
std::this_thread::sleep_for(fiveMillis);

// prints an implementation-defined thread handle
std::cout << std::this_thread::get_id() << std::endl;

// thread-local variable - useful only for non-local variables
thread_local int counter = 0;

void f()
{
    int counter = 0;          // automatically per thread
}

std::this_thread::yield();    // allow other threads to run
```

- `thread_local` is a portable version of `pthread_getspecific()` and `TlsGetValue()`

Mutexes

```
// C++11
// may block if locked more than once by the same thread
std::mutex m;
m.lock();
...
m.unlock();

if (m.try_lock()) { // test if mutex can be acquired
    std::cout << "I got the lock!\n";
    m.unlock();
}

// wait up to a certain amount of time to acquire a lock
auto locked = m.wait_for(std::chrono::seconds(10));

auto timeT = std::chrono::steady_clock::now()
              + std::chrono::seconds(10);
auto locked2 = m.wait_until(timeT);

// recursive mutex can be locked more than once by the same thread
std::recursive_mutex rm;
```

Mutexes and exception safety

```
// C++11
std::mutex m;

void print(int i)
{
    std::lock_guard<std::mutex> lock(m);
    f(i);          // might throw
    std::cout << "i = " << i << std::endl;
}
```

- `std::lock_guard<>` guarantees that `m` will be unlocked even if an exception is thrown

Condition variables

```
std::condition_variable start;
std::mutex mutex;
std::unique_lock<std::mutex>
    startMutex(mutex);
int dataItems = 0;

void workerCalc()
{
    startMutex.lock();
    start.wait(startMutex,
        []{ return dataItems > 0; });
    claimDataItem();
    startMutex.unlock();
    workOnDataItem();
}
```

```
void setUpWorkers()
{
    startMutex.lock();
    dataItems = N;
    startMutex.unlock();
    start.notify_all();
}
```

blocks until
condition is
true

signal all threads
waiting for
condition

- How to wait until some logical condition is true
 - an alternative to polling
- Need to acquire a lock then wait on condition
 - lock is unlocked during the wait and relocked on wakeup
 - internal loop testing predicate to avoid “spurious wakeups”

One-time initialisation

```
std::once_flag onceFlag;
std::ofstream logFile;
std::mutex logMutex;

void initLogFunc()
{
    system("rm -rf /opt/myapp/output/");
    logFile.open("/opt/myapp/myapp.log", ios::ate | ios::app);
    if (! logFile) /* handle error */
}

void log(const std::string & msg)
{
    std::call_once(onceFlag, initLogFunc);
    std::lock_guard<std::mutex> guard(logMutex);
    using namespace std::chrono;
    log << system_clock::to_time_t(system_clock::now())
        << "\t" << msg << std::endl;
}
```

- `std::once_flag` provides a means to guarantee that only one thread will run a piece of code only once

Threads and locks

- Up to now we have not used explicit locks so there has been no possibility of deadlock
- With explicit thread-and-lock programming we enter a whole new world of potential pain
 - use parallel libraries, futures and atomics if you can!
 - if you can't, use only a single lock if performance allows
- Handling multiple locks
 - incorrect locking order can lead to deadlock
 - how to move, adopt or defer locking

std::lock

```
std::mutex lock1, lock2;  
  
std::lock(lock1, lock2); // guarantees all locks acquired or none  
  
std::lock_guard<std::mutex> guard1(lock1, std::adopt_lock);  
std::lock_guard<std::mutex> guard2(lock2, std::adopt_lock);
```

- Which order should we lock lock1 and lock2?
 - If different threads do it in different orders then deadlock
- std::lock provides a non-deadlocking ordering and releases any locks it doesn't acquire (or if an exception is thrown)
 - handles multiple locks (variadic function template)
- Hand over already locked lock to std::lock_guard for exception-safe unlocking

std::unique_lock

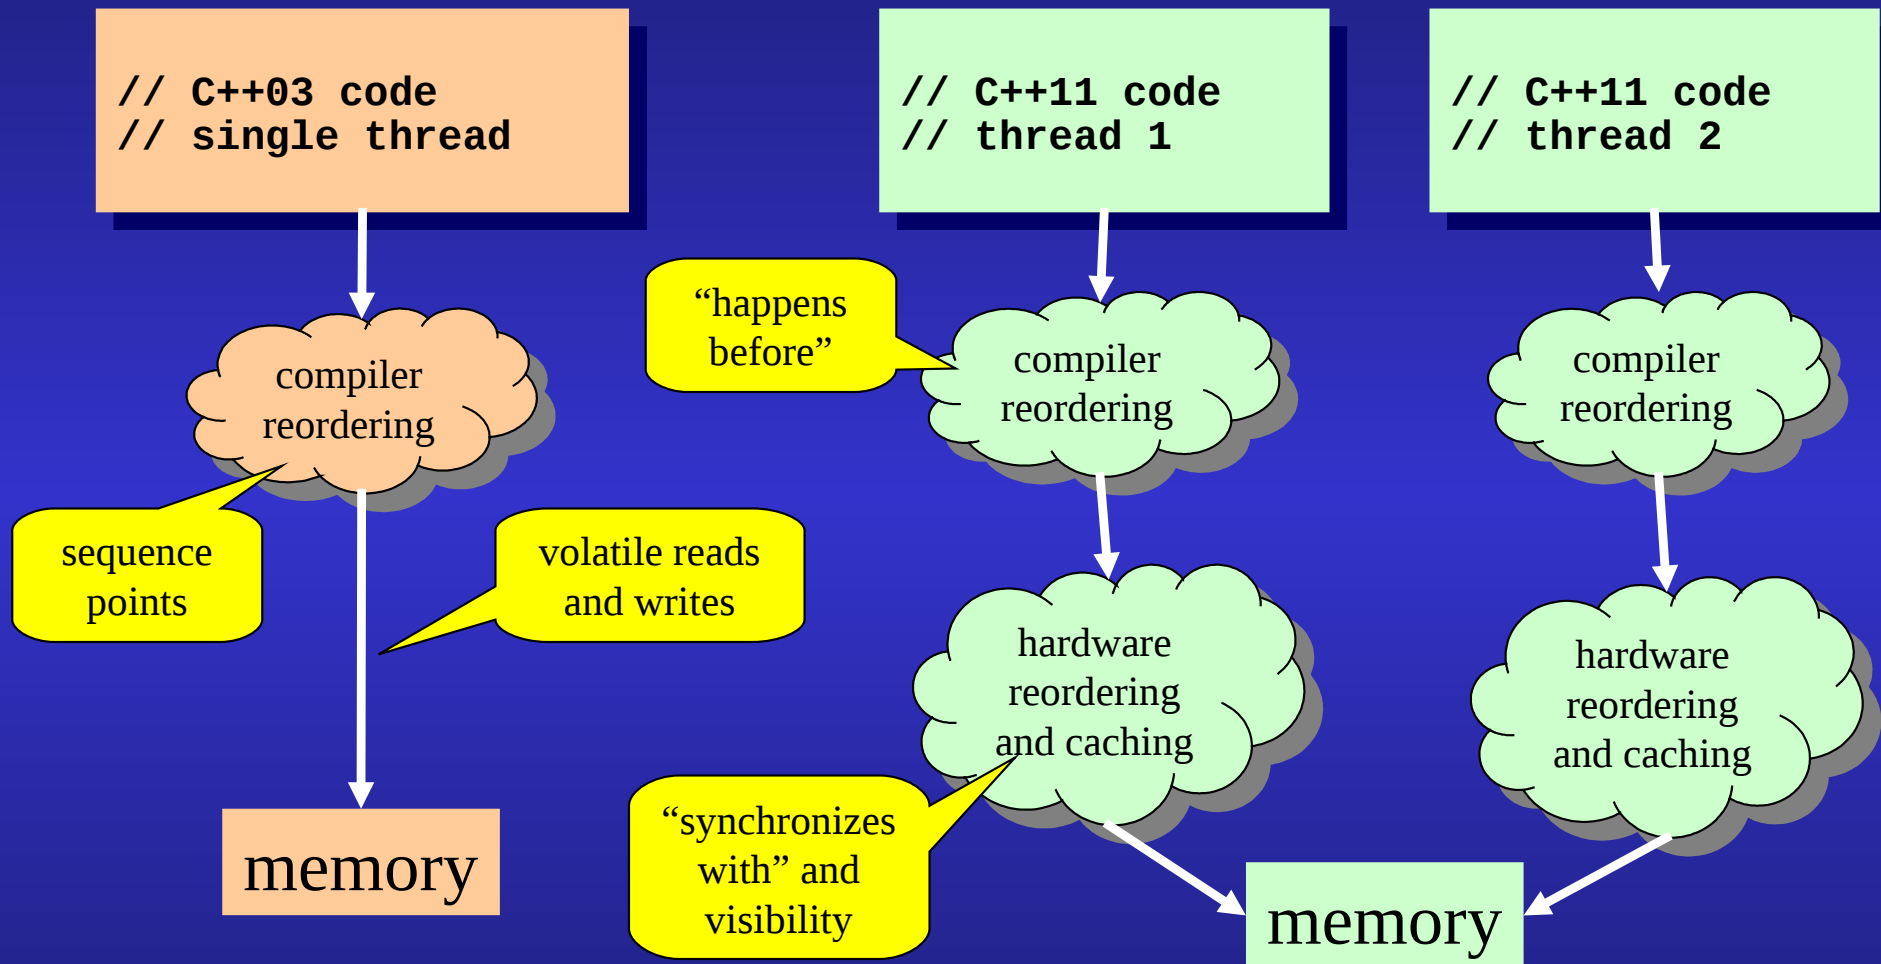
```
std::mutex lock1, lock2;  
  
std::unique_lock<std::mutex> uniq1(lock1, std::defer_lock);  
std::unique_lock<std::mutex> uniq2(lock2, std::defer_lock);  
  
std::lock(lock1, lock2); // guarantees all locks acquired or none
```

- std::unique_lock<> allows for moving locks, adopting already acquired locks and referring to locks without acquiring them
 - more flexible version of std::lock_guard
 - small performance overhead
 - can also adopt locks

Memory model

- C++11, unlike C++03, embraces the notion of more than one thread of execution
- In order to support this, C++11 had to create a memory model to define what operations are visible in other threads and when
- Needed to control optimisations by the compiler and the hardware
 - visibility and reordering of operations
- Lack of standards has meant such code is not portable up till now

Need for a memory model



- Correctness now based on memory, not just code
- Need to control caching (register, L1, L2, etc)

Instruction interleaving

```
// thread 1  
x = 1;  // 1  
r1 = y; // 2
```

```
// thread 2  
y = 1;  // 3  
r2 = x; // 4
```

interleave



```
// 6 possible sequentially consistent  
// execution orders
```

```
1234 // x=1; r1=y; y=1; r2=x;  
1324 // x=1; y=1; r1=y; r2=x;  
1342 // x=1; y=1; r2=x; r1=y;  
3412 // y=1; r2=x; x=1; r1=y;  
3142 // y=1; x=1; r2=x; r1=y;  
3124 // y=1; x=1; r1=y; r2=x;
```

- “Sequential consistency” means operations in separate threads are interleaved and that all threads see the same interleaving
 - Sequence is preserved within and across threads
- This is the “natural” mental model for programmers to think of thread execution order and memory
 - It is also the C++11 default memory ordering

Instruction reordering

```
// thread 1  
x = 1;  // 1  
r1 = y; // 2
```

```
// thread 2  
y = 1;  // 3  
r2 = x; // 4
```

reorder



```
// 4 factorial (== 24)  
// possible execution orders
```

```
1234 // x=1; r1=y; y=1; r2=x;  
4321 // r2=x; y=1; r1=y; x=1;  
3124 // y=1; x=1; r1=y; r2=x;  
// etc...
```

could be
executed in
reverse order

- In order to gain performance both the compiler and the hardware may reorder instructions
 - compiler may move loads earlier (to allow for cache misses)
 - hardware may not write back to memory immediately (store buffers)
- x, y, r1 and r2 are all independent so code can be reordered
- Even worse, changes in one thread may not be visible in another thread so results are not defined – **data race**

Hardware memory reordering

	Alpha	ARMv7	PA-RISC	POWER	SPARC RMO	SPARC PSO	SPARC TSO	x86	x86 oostore	AMD64	IA-64	zSeries
Loads reordered after loads	Y	Y	Y	Y	Y				Y		Y	
Loads reordered after stores	Y	Y	Y	Y	Y				Y		Y	
Stores reordered after stores	Y	Y	Y	Y	Y	Y			Y		Y	
Stores reordered after loads	Y	Y	Y	Y	Y	Y	Y	Y	Y	Y	Y	Y
Atomic reordered with loads	Y	Y		Y	Y						Y	
Atomic reordered with stores	Y	Y		Y	Y	Y					Y	
Dependent loads reordered	Y											
Incoherent Instruction cache pipeline	Y	Y		Y	Y	Y	Y	Y	Y		Y	Y

http://en.wikipedia.org/wiki/Memory_ordering

- Hardware can reorder memory operations in different ways
 - Can also depend on operating system
 - Solaris on SPARC uses Total Store Order (TSO)
 - Linux on SPARC uses Relaxed Memory Order (RMO)

Synchronisation with seq. consistency

```
// thread 1  
x = 42;  
x_init = true;
```

```
// thread 2  
while (! x_init) {}  
y = x;
```

- This code is correct when sequentially consistent
 - thread 2 doesn't access x until it has been set by thread 1
- But in the presence of reordering it can fail
- The problem is that we haven't specified that cross-thread order or visibility is important
- We need to use synchronisation variables – atomics
- Making everything atomic is slow – 30-60 cycles
 - cache synchronisation is slow and has limited bandwidth

Synchronisation with atomics

```
// thread 1
std::atomic<bool> x_init;
int x;

x = 42;
x_init.store(true);

// or x_init = true;
```

```
// thread 2
extern std::atomic<bool> x_init;
extern int x;

while (! x_init.load());
y = x;

// or while (! x_init);
```

- This now works without relying on having sequential consistency everywhere (just atomics)
 - atomics prevent the compiler moving code across accesses
 - atomics also cause memory updates to be visible
- Load and store uses sequential consistency
 - uses default parameter of `std::memory_order_seq_cst`

Low-level synchronisation detail

```
// thread 1
std::atomic<bool> x_init;
```

```
x = 42;
// blue fence
x_init.store(true);
// red fence
```

prevents x and x_init
being reordered

makes store to
x_init visible

```
// thread 2
extern std::atomic<bool> x_init;
```

```
while (/*rfence*/ ! x_init.load());
// blue fence
y = x;
```

prevents
reordering

rfence here in
loop forces load
of latest value
of x_init

- Blue fences prevent the compiler reordering code
 - they don't generate any run-time code
- Red fences force memory to make changes visible
 - they do generate code: fence, lock prefix, CAS opcodes
 - depends heavily on underlying hardware (c.f. reordering)
 - only need one of the two red fences, usually on store
- One reason that threads can't just be a library

Generated assembler code

```
// thread 1
x = 42;
// blue fence
x_init.store(true);
// red fence
```

```
// with atomic bool x_init
mov     DWORD PTR x, 42
mov     BYTE PTR x_init, 1
mfence

// with bool x_init
mov     DWORD PTR x,
mov     BYTE PTR x_init, 1
```

g++ 4.7
output

fence needed
on store for
X86

infinite loop
because
visibility not
specified

```
// thread 2

while (/*rfence*/ ! x_init.load());
// blue fence
y = x;
```

sequential consistent
by default

```
// with atomic bool x_init
.L25:
movzx   eax, BYTE PTR x_init
test    al, al
je      .L25
mov     eax, DWORD PTR x
mov     DWORD PTR y, eax

// with bool x_init
cmp     BYTE PTR x_init, 0
jne     .L3
.L5:
jmp     .L5
.L3:
mov     eax, DWORD PTR x
mov     DWORD PTR y, eax
```

no fence
needed on load
for X86

Using memory order flags

```
// thread 1
x = 42;
// blue fence
x_init.store(true,
std::memory_order_release);
```

```
// thread 2
while (!
x_init.load(std::memory_order_acquire));
// blue fence
y = x;
```

acquire means no
reads in this thread
reordered before here

```
// with bool x_init seq_cst
mov     DWORD PTR x, 42
mov     BYTE PTR x_init, 1
mfence
```

needed on
seq_cst store

```
// with bool x_init release
mov     DWORD PTR x, 42
mov     BYTE PTR x_init, 1
```

no fence for
release store

```
// with bool atomic x_init
L25:
movzx   eax, BYTE PTR x_init
test    al, al
je       .L25
mov     eax, DWORD PTR x
mov     DWORD PTR y, eax

// same code for x_init acquire
```

no fence
needed on load
for X86

- Release provides only the blue fence (no writes in this thread reordered after the store)
- Controls compiler reordering but not hardware

Memory model advice

- This is a complex and subtle area and you should avoid using it unless you can prove that you can't get adequate performance without it
 - yes, really, I mean it....
- Even experts get confused by this stuff!
 - did I mention you should avoid it....
- If you do use it, use acquire on load and release on store
 - Anything else will be a source of subtle bugs