

Use a more efficient algorithm than the copy-and-sort version in the slides.

```
#include <type_traits>
#include <iostream>
#include <utility>

template <typename T>
T max2(T && t)
{
    return t;
}

template <typename T, typename... Args>
typename std::common_type<T, Args...>::type
max2(T && t, Args&& ... args)
{
    typedef typename std::common_type<T, Args...>::type C;
    return std::max(t, max2(std::forward<C>(args)...));
}

int main10()
{
    std::cout << max2(5, 14u, 22.1, 9.3, 13) << std::endl;
    return 0;
}
```

Basic concurrency

- 1) Using `std::async`, open multiple files in parallel and count the total number of characters across all the files

```
#include <iostream>
#include <fstream>
#include <thread>
#include <future>
#include <string>
#include <algorithm>
#include <iterator>
#include <vector>

std::size_t countCharsInFile(const std::string & fileName)
{
    std::ifstream input(fileName);
    if (! input)
        std::cout << "Can't open " << fileName << std::endl;
    return std::distance(std::istreambuf_iterator<char>(input),
                        std::istreambuf_iterator<char>());
}

int main(int argc, char * argv[])
{
    std::vector<std::future<std::size_t>> futures;
    for (auto i = 1; i != argc; ++i) {
        auto f = std::async(std::launch::async, countCharsInFile, argv[i]);
        futures.push_back(std::move(f));
    }
}
```

```

    std::size_t total = 0;
    for (auto & f : futures)
        total += f.get();

    std::cout << "Total characters = " << total << std::endl;
    return 0;
}

```

- 2) Fill a vector with random integers in the range 1 to 1000. Count the number of odd integers using the optimum number of threads

```

#include <random>
#include <vector>
#include <thread>
#include <algorithm>
#include <iostream>
#include <numeric>
#include <iterator>
#include <future>

int sum(std::vector<int>::const_iterator first,
std::vector<int>::const_iterator last)
{
    return std::accumulate(first, last, 0);
}

int main2()
{
    using namespace std;
    default_random_engine engine;
    uniform_int_distribution<> ints(1, 1000);

    std::vector<int> v;
    const auto numElems = 5000u;
    const auto minGrain = 1000u;
    const auto numThreads = min(numElems / minGrain,
                                thread::hardware_concurrency());
    generate_n(back_inserter(v), numElems, [&]{ return ints(engine); });
    const auto chunk = numElems / numThreads;
    std::vector<future<int>> futures;
    for (auto i = 0; i != numThreads; ++i) {
        const auto begin = v.cbegin()+i*chunk;
        const auto end = v.cbegin()+min((i+1)*chunk, v.size());
        auto f = async(launch::async, sum, begin, end);
        futures.push_back(std::move(f));
    }
    int total = 0;
    for (auto & f : futures)
        total += f.get();
    cout << "Total = " << total << endl;
    return 0;
}

```

3) Implement a simple spin lock using an atomic

```
#include <atomic>
#include <future>
#include <vector>
#include <iostream>

std::atomic<bool> spinlock(false);

void lock_unlock()
{
    for (auto i = 0; i != 10'000; ++i) {
        bool expect = false;
        while (! spinlock.compare_exchange_weak(expect, true)) {
            expect = false;
        }
        std::cout << "got ";
        std::cout << "the ";
        std::cout << "lock\n";
        spinlock = false;
    }
}

int main3()
{
    const int numThreads = 10;

    std::vector<std::future<void>> futures;
    for (auto i = 0; i != numThreads; ++i)
        futures.emplace_back(std::async(std::launch::async, lock_unlock));
}
```

4) Initialise a global shared object pointed to by an atomic pointer and make sure the object is correctly deleted when the program exits. Do this without using any locks

```
#include <atomic>
#include <iostream>
#include <future>
#include <vector>

struct X { int stuff[100]; };

std::atomic<X*> xp{nullptr};

void initxp()
{
    if (xp == nullptr) {
        X * newX = new X;
        X * expected = nullptr;
        auto updated = xp.compare_exchange_strong(expected, newX);
        if (! updated) {
            std::cout << "didn't get it\n";
            delete newX;
        } else {
            std::cout << "got it\n";
        }
    } else {
        std::cout << "already assigned\n";
    }
}
```

```

    }
}

int main4()
{
    {
        const int numThreads = 10;
        std::vector<std::future<void>> futures;
        futures.reserve(numThreads);
        for (auto i = 0; i != numThreads; ++i)
            futures.emplace_back(std::async(std::launch::async, initxp));
    }
    delete xp;
}

```

- 5) Insert items at the head of a singly-linked list from multiple threads using a non-blocking algorithm using compare and exchange (CAS). Check that the correct number of items have been inserted into the list and that the list structure is correct

```

#include <iostream>
#include <atomic>
#include <thread>
#include <vector>

const auto numLoopsPerThread = 10 * 1000;
const auto numThreads = 4;

struct Node {
    Node * next;
    std::thread::id who;
};

std::atomic<Node *> head(nullptr);

void addToList()
{
    for (auto i = 0; i != numLoopsPerThread; ++i) {
        Node * newNode = new Node;
        newNode->next = head;
        newNode->who = std::this_thread::get_id();
#ifdef 0
        head = newNode; // not thread safe!
#else
        while (! head.compare_exchange_weak(newNode->next, newNode))
            std::cout << "looping on head\n";
#endif
    }
}

int main5()
{
    std::vector<std::thread> threads;
    for (auto i = 0; i != numThreads; ++i)
        threads.emplace_back(addToList);
    for (auto & t : threads)
        t.join();

    int nodeCount = 0;
    while (head.load()) {

```

```

        nodeCount++;
        Node * p = head;
        head = head->next;
        delete p;
    }

    std::cout << "Node count = " << nodeCount << std::endl;

    return 0;
}

```

- 6) Create a single-producer multi-consumer queue based on a circular buffer. Use atomics to hold the positions for the writer and the readers. Use non-blocking queue operations throughout (such as spin locks or polling loops)

```

#include <iostream>
#include <atomic>
#include <vector>
#include <thread>
#include <typeinfo>
#include <limits>
#include <utility>

//const auto numItemsToWrite = 30 * 1000;
const auto numItemsToWrite = 20;
//const auto bufSize = 10 * 1000;
const auto bufSize = 10;
const auto numReaders = 4;

struct Item { int value; };

Item buf[bufSize];

std::atomic<int> readerPosns[numReaders] = { 0 };
std::atomic<int> writerPosn = 0;

int minReaderPosn()
{
    int minPosn = std::numeric_limits<int>::max();
    for (auto i = 0; i != numReaders; ++i)
        minPosn = std::min(minPosn, readerPosns[i].load());
    return minPosn;
}

class Writer {
public:
    Writer() : posn(0) {}
    void fillBuffer() {
        for (auto i = 0; i != numItemsToWrite; ++i) {
            while (posn >= minReaderPosn() + bufSize) {
                std::cout << "Waiting to write" << std::endl;
                std::this_thread::sleep_for(std::chrono::milliseconds(1));
            }
            buf[posn % bufSize].value = posn;
            std::cout << "Wrote item " << posn << std::endl;
            posn++;
            writerPosn++;
        }
    }
}

```

```

    }

private:
    int posn;
};

class Reader {
public:
    Reader(int id) : id(id), posn(0) {}
    void readBuffer() {
        for (auto i = 0; i != numItemsToWrite; ++i) {
            while (posn >= writerPosn) {
                std::cout << "Reader " << id << " waiting to read"
                    << std::endl;
                std::this_thread::sleep_for(std::chrono::milliseconds(1));
            }
            std::cout << "Reader " << id << " read item " << posn
                << std::endl;
            readerPosns[id]++;
            posn++;
        }
    }

private:
    int id, posn;
};

int main6()
{
    std::vector<std::thread> readers;
    std::thread writer(&Writer::fillBuffer, Writer());
    for (auto i = 0; i != numReaders; ++i)
        readers.emplace_back(&Reader::readBuffer, Reader(i));
    for (auto & t : readers)
        t.join();
    writer.join();
    return 0;
}

```

More concurrency

- 1) Update the RUList to be thread-safe by locking the appropriate methods. Is this a sensible design approach?

```

#include <iostream>
#include <mutex>
#include <vector>
#include <string>
#include <algorithm>

template <typename T>
class RUList
{
public:
    RUList() = default;
    RUList(std::initializer_list<T> lst) : v(lst) {
        if (v.size() > maxSize)

```

```

        v.resize(maxSize);
    }

    RUList(const RUList & other) {
        std::lock_guard<std::mutex> other_guard(other.m);
        v = other.v;
    }

    RUList(RUList && other) {
        std::lock_guard<std::mutex> other_guard(other.m);
        v = std::move(other.v);
    }

    RUList & operator=(const RUList & rhs) {
        RUList temp(rhs);
        std::lock(m, temp.m);
        std::lock_guard<std::mutex> guard_m(m, std::adopt_lock);
        std::lock_guard<std::mutex> guard_t(temp.m, std::adopt_lock);
        std::swap(v, temp.v);
        return *this;
    }

    RUList & operator=(RUList && rhs) {
        std::lock(m, rhs.m);
        std::lock_guard<std::mutex> guard_m(m, std::adopt_lock);
        std::lock_guard<std::mutex> guard_r(rhs.m, std::adopt_lock);
        v = std::move(rhs.v);
        return *this;
    }

    template <typename T2>
    void add(T2 && name) {
        std::lock_guard<std::mutex> guard(m);
        auto i = std::find(begin(v), end(v), std::forward<T2>(name));
        if (i != end(v))
            v.erase(i);
        v.insert(v.begin(), std::forward<T2>(name));
        if (v.size() > maxSize)
            v.pop_back();
        print();
    }

    void print() const {
        std::lock_guard<std::mutex> guard(m);
        std::cout << "-----\n";
        for (const auto & i : v)
            std::cout << i << "\n";
    }

    const T & operator[](unsigned int i) const noexcept {
        std::lock_guard<std::mutex> guard(m);
        return v[i];
    }

private:
    std::vector<T> v;
    mutable std::mutex m;
    static constexpr unsigned int maxSize = 3;
};

RUList<std::string> f() { return RUList<std::string>{"one elem"}; }

```

```

int main()
{
    RUList<std::string> l{"a", "b", "c", "d"};
    l.add("foo");
    l.add("foo");
    l.add("bar");
    l.add("bar");
    l.add("wibble");
    l.add("xyzzzy");
    l.add("wibble");

    RUList<std::string> l2(std::move(l));
    l.print();
    l2.add("123");

    RUList<std::string> x;
    x = f();
    std::cout << x[0] << std::endl;

    RUList<std::string> l3(l2);
    l2 = l3;
}

```

- 2) Create two bank accounts, each with its own associated lock. Transfer random amounts based on a normal distribution between the two accounts in both directions using several parallel threads. Ensure that the accounts are locked and unlocked correctly to avoid deadlock
- 3) Initialise the random number engine in the previous exercise using a hardware random number source from only one of the threads

```

#include <iostream>
#include <mutex>
#include <thread>
#include <random>
#include <vector>
#include <algorithm>

class Account {
public:
    Account() : balance(0.0) {}
    void deposit(double amount) { balance += amount; }
    void withdraw(double amount) { balance -= amount; }
    double getBalance() const { return balance; }
    std::mutex & getLock() { return m; }
    void lock() { m.lock(); }
    void unlock() { m.unlock(); }
private:
    double balance /* = 0.0 */;
    std::mutex m;
};

std::default_random_engine engine;
std::normal_distribution<> amounts(0.0, 100.0);
std::mutex dreMutex;
std::once_flag callOnceFlag;

```



```

void setup()
{
    std::random_device rd;
    engine.seed(rd());
}

void transfer(Account & a1, Account & a2)
{
    std::call_once(callOnceFlag, setup);

    const auto numTransfers = 1000;
    for (auto i = 0; i != numTransfers; ++i) {
        dreMutex.lock();
        auto amount = amounts(engine);
        dreMutex.unlock();
        std::lock(a1.getLock(), a2.getLock());
        std::lock_guard<std::mutex> guarda1(a1.getLock(), std::adopt_lock);
        std::lock_guard<std::mutex> guarda2(a2.getLock(), std::adopt_lock);
        if (amount > 0.0) {
            a1.withdraw(amount);
            a2.deposit(amount);
        } else if (amount < 0.0) {
            a2.withdraw(-amount);
            a1.deposit(-amount);
        }
    }
}

int main4()
{
    Account a1, a2;

    std::vector<std::thread> tthrds;
    const auto numThreads = 10;
    for (auto i = 0; i != numThreads; ++i)
        tthrds.emplace_back(transfer, std::ref(a1), std::ref(a2));
    std::for_each(tthrds.begin(), tthrds.end(),
        std::mem_fn(&std::thread::join));
    std::cout << a1.getBalance() << " " << a2.getBalance() << std::endl;
    return 0;
}

```

- 4) Implement a thread-safe fixed-length queue using blocking operations (instead of spin locks or polling)

```

#include <iostream>
#include <deque>
#include <thread>
#include <mutex>
#include <condition_variable>

template <typename T>
class Queue {
public:
    void push(const T & t) {
        std::unique_lock<std::mutex> fullLock(queueMutex);
        full.wait(fullLock, [this]{ return queue.size() < maxQueueLen; });
        queue.push_back(t);
        std::cout << "Pushed item " << t << std::endl;
    }
};

```

```

        empty.notify_one();
    }

    void pop(T & t) {
        std::unique_lock<std::mutex> emptyLock(queueMutex);
        empty.wait(emptyLock, [this]{ return ! queue.empty(); });
        t = std::move(queue.front());
        queue.pop_front();
        std::cout << " Popped item " << t << std::endl;
        full.notify_one();
    }

private:
    std::deque<T> queue;
    std::mutex queueMutex;
    std::condition_variable empty, full;
    static const int maxQueueLen = 4;
};

Queue<int> queue;

const auto numItemsToInsert = 30;

void writer()
{
    for (auto i = 0; i != numItemsToInsert; ++i)
        queue.push(i+1);
}

void reader()
{
    int item;
    for (auto i = 0; i != numItemsToInsert; ++i)
        queue.pop(item);
}

int main5()
{
    std::thread writeThread(writer);
    std::thread readThread(reader);
    writeThread.join();
    readThread.join();
    return 0;
}

```

- 5) Add push and pop methods that will time out if they block for longer than a given time.
Refactor out any duplication

```

#include <iostream>
#include <deque>
#include <thread>
#include <mutex>
#include <condition_variable>
#include <chrono>

```

```

template <typename T>
class Queue {
public:
    Queue() :
        emptyPred([this]{ return ! queue.empty(); }),
        fullPred([this]{ return queue.size() < maxQueueLen; }) {}

    void push(const T & t) {
        std::unique_lock<std::mutex> fullLock(queueMutex);
        full.wait(fullLock, fullPred);
        queue.push_back(t);
        std::cout << "Pushed item " << t << std::endl;
        empty.notify_one();
    }

    void push_wait_for(const T & t, const std::chrono::milliseconds & dur) {
        std::unique_lock<std::mutex> fullLock(queueMutex);
        auto ok = full.wait_for(fullLock, dur, fullPred);
        if (ok) {
            queue.push_back(t);
            std::cout << "Pushed item " << t << std::endl;
            empty.notify_one();
        } else {
            std::cout << "*** Waited too long\n";
        }
    }

    void pop(T & t) {
        std::unique_lock<std::mutex> emptyLock(queueMutex);
        empty.wait(emptyLock, emptyPred);
        t = std::move(queue.front());
        queue.pop_front();
        std::cout << "Popped item " << t << std::endl;
        full.notify_one();
    }

private:
    std::deque<T> queue;
    std::mutex queueMutex;
    std::condition_variable empty, full;
    std::function<bool()> emptyPred, fullPred;
    static const int maxQueueLen = 4;
};

Queue<int> queue;

const auto numItemsToInsert = 30;

void writer()
{
    for (auto i = 0; i != numItemsToInsert; ++i)
        queue.push_wait_for(i+1, std::chrono::milliseconds(1));
}

void reader()
{
    int item;
    for (auto i = 0; i != numItemsToInsert; ++i)
        queue.pop(item);
}

```

```

int main6()
{
    std::thread writeThread(writer);
    writeThread.join();
    return 0;
}

```

6) Add appropriate memory ordering flags to the circular buffer queue from the previous section

```

#include <iostream>
#include <atomic>
#include <vector>
#include <thread>
#include <typeinfo>
#include <limits>
#include <utility>

const auto numItemsToWrite = 20;
const auto bufSize = 10;
const auto numReaders = 4;

struct Item { int value; };

Item buf[bufSize];

std::atomic<int> readerPosns[numReaders];
std::atomic<int> writerPosn;

int minReaderPosn()
{
    int minPosn = std::numeric_limits<int>::max();
    for (auto i = 0; i != numReaders; ++i)
        minPosn = std::min(minPosn,
                           readerPosns[i].load(std::memory_order_relaxed));
    return minPosn;
}

class Writer {
public:
    Writer() : posn(0) {}
    void fillBuffer() {
        for (auto i = 0; i != numItemsToWrite; ++i) {
            while (posn >= minReaderPosn() + bufSize) {
                std::cout << "Waiting to write" << std::endl;
                std::this_thread::sleep_for(std::chrono::milliseconds(1));
            }
            buf[posn % bufSize].value = posn;
            std::cout << "Wrote item " << posn << std::endl;
            posn++;
            writerPosn.fetch_add(1, std::memory_order_release);
        }
    }
private:
    int posn;
};

```

```

class Reader {
public:
    Reader(int id) : id(id), posn(0) {}
    void readBuffer() {
        for (auto i = 0; i != numItemsToWrite; ++i) {
            while (posn >= writerPosn.load(std::memory_order_acquire)) {
                std::cout << "Reader " << id << " waiting to read"
                    << std::endl;
                std::this_thread::sleep_for(std::chrono::milliseconds(1));
            }
            std::cout << "Reader " << id << " read item " << posn
                << std::endl;
            posn++;
            readerPosns[id].fetch_add(1, std::memory_order_release);
        }
    }

private:
    int id, posn;
};

int main7()
{
    std::fill(std::begin(readerPosns), std::end(readerPosns), 0);
    writerPosn = 0;
    std::vector<std::thread> readers;
    std::thread writer(&Writer::fillBuffer, Writer());
    for (auto i = 0; i != numReaders; ++i)
        readers.emplace_back(&Reader::readBuffer, Reader(i));
    for (auto & t : readers)
        t.join();
    writer.join();
    return 0;
}

```