

Use a more efficient algorithm than the copy-and-sort version in the slides.

### Basic concurrency

- 1) Using `std::async`, open multiple files in parallel and count the total number of characters across all the files
- 2) Fill a vector with random integers in the range 1 to 1000. Count the number of odd integers using the optimum number of threads
- 3) Implement a simple spin lock using an atomic
- 4) Initialise a global shared object pointed to by an atomic pointer and make sure the object is correctly deleted when the program exits. Do this without using any locks
- 5) Insert items at the head of a singly-linked list from multiple threads using a non-blocking algorithm using compare and exchange (CAS). Check that the correct number of items have been inserted into the list and that the list structure is correct
- 6) Create a single-producer multi-consumer queue based on a circular buffer. Use atomics to hold the positions for the writer and the readers. Use non-blocking queue operations throughout (such as spin locks or polling loops)

### More concurrency

- 1) Update the RUList to be thread-safe by locking the appropriate methods. Is this a sensible design approach?
- 2) Create two bank accounts, each with its own associated lock. Transfer random amounts based on a normal distribution between the two accounts in both directions using several parallel threads. Ensure that the accounts are locked and unlocked correctly to avoid deadlock
- 3) Initialise the random number engine in the previous exercise using a hardware random number source from only one of the threads
- 4) Implement a thread-safe fixed-length queue using blocking operations (instead of spin locks or polling)
- 5) Add push and pop methods that will time out if they block for longer than a given time. Refactor out any duplication
- 6) Add appropriate memory ordering flags to the circular buffer queue from the previous section